

Lecture 8

*Lecturer: Andre Wibisono**Scribe: Anay Mehrotra*

1 Outline

Today's lecture covers two main topics:

1. Effective diameter and concentration inequalities for strongly log-concave distributions, and
2. lower bounds on the s -conductance of Markov chains.

2 Preliminaries

Notation. Given $x \in \mathbb{R}^n$ and $r \geq 0$, let $\mathbb{B}(x, r) \subseteq \mathbb{R}^n$ denote the ℓ_2 -ball centered at x with radius r .¹ Given two distributions ν and ρ on \mathbb{R}^n , define the total variation distance between them as

$$\text{TV}(\nu, \rho) := \sup_{A \subseteq \mathbb{R}^n} |\nu(A) - \rho(A)|.$$

Given two sets $S, T \subseteq \mathbb{R}^n$, define the distance between them to be

$$d(S, T) := \inf_{x \in S, y \in T} \|x - y\|_2.$$

2.1 Markov chains

Let P be a Markov chain. Suppose that the current point of P is $x \in \mathbb{R}^n$. Let P_x denote the density of the next point of P . For a Markov chain P , define the ergodic flow from a set $A \subseteq \mathbb{R}^n$ to A^c as

$$\Phi(A) := \int_A P_x(A^c) \, d\nu(x). \quad (\text{Ergodic flow})$$

In other words, $\Phi(A)$ is the probability that P moves from a point in A (sampled from ν) to a point in A^c .

Definition 1 (s -conductance). For $s \in [0, 1/2)$ define the s -conductance of P as

$$\Phi_s(P) := \inf_{A \subseteq \mathbb{R}^n: s < \nu(A) < 1-s} \frac{\Phi(A)}{\min\{\nu(A) - s, \nu(A^c) - s\}}. \quad (s\text{-conductance})$$

¹In other words, $B(x, r) := \{y \in \mathbb{R}^n: \|y - x\|_2 \leq r\}$.

2.2 Strong log-concavity and strong convexity

From the last lecture, recall that for any $\alpha > 0$ a twice differentiable function, f , is said to be α -strongly convex if $\nabla^2 f(x) \succeq \alpha I$. Equivalently, f is said to be α -strongly convex if the smallest eigenvalue of $\nabla^2 f(x)$ is at least α . Further, a distribution ν is said to be α -strongly logconcave if the function $f: \mathbb{R}^n \rightarrow \mathbb{R}$ defined as $f := -\log(\nu)$ is α -strongly convex. In particular, this implies that a α -strongly convex distribution has a unique mode. For further discussion of log-concave distributions, we refer the reader to [Vem05, Section 2] and [LV07, Section 5].

Example 1. Consider the Gaussian distribution $\rho = \mathcal{N}(\nu, \Sigma)$ with mean $\nu \in \mathbb{R}^n$ covariance $\Sigma \in \mathbb{R}^{n \times n}$:

$$\rho(x) := \frac{1}{\sqrt{\det(2\pi\Sigma)}} \cdot \exp\left(-\frac{1}{2} \cdot (x - \mu)^\top \Sigma^{-1} (x - \mu)\right).$$

In this example, $f(x) := -\log(\rho) \circ (x) = \frac{1}{2} \cdot (x - \mu)^\top \Sigma^{-1} (x - \mu) + \frac{1}{2} \cdot \log \det(2\pi\Sigma)$. The hessian of f is $\nabla^2 f(x) = \Sigma^{-1}$. Therefore, f is $\frac{1}{\lambda_{\max}(\Sigma)}$ -strongly convex and ρ is $\frac{1}{\lambda_{\max}(\Sigma)}$ -strongly log-concave.

2.3 Isoperimetry for strongly logconcave distributions

Recall the following isoperimetry-inequality for strongly logconcave distributions presented in Lecture 7.

Theorem 1 (Isoperimetry constant for strongly logconcave distributions). All α -strongly log-concave distribution ν on \mathbb{R}^n satisfy isoperimetry with constant

$$\Psi \geq \log 2 \cdot \sqrt{\alpha}.$$

That is, for any partition $S \cup T \cup (\mathbb{R}^n \setminus S \setminus T)$, it holds that

$$\nu(\mathbb{R}^n \setminus S \setminus T) \geq \log 2 \cdot \sqrt{\alpha} \cdot d(S, T) \cdot \min\{\nu(S), \nu(T)\}. \quad (1)$$

Interestingly, the isoperimetry constant Ψ is independent of the dimension n for strongly logconcave distributions. Kannan, Lovász and Simonovits conjectured that this is true for all log-concave distributions (not just strongly log-concave ones); this is known as the KLS conjecture [KLS95]. Since its proposal in 1995, it has been a subject of intense study—both in attempts to resolve the conjecture [Kla07b, FGP07, Kla07a, Fle10, GM11, LV17a] and in establishing its relation to other well-known conjectures [Bal06, Eld13]. We refer the reader to the excellent survey [LV18] for an in-depth discussion and several more references.

Typically proofs of lower bounds on isoperimetry constants proceed through the idea of “Localization,” which reduces inequalities in measures of sets (i.e., integrals in \mathbb{R}^n) to one-dimensional integrals. At a high-level, the reduction proceeds by starting with \mathbb{R}^n and repeatedly bisecting the space so that the original inequality holds in the remaining space after the bisection. In the limit, this process reduces the dimension of the space from n to 1. We refer the reader to [YTL19, Section 9.5] for an overview of localization; and [Eld13, EL14, Fle10, LV17b] for a discussion of state-of-the-art localization techniques.

3 Effective diameter and concentration

In this section, we motivate the notion of effect diameter and present concentration result for strongly log-concave distributions. Consider an α -strongly log-concave distribution, $\nu \propto e^{-f}$ on \mathbb{R}^n . Let \bar{x} be the mean of ν , i.e., $\bar{x} := \frac{\int_{\mathbb{R}^n} x e^{-f(x)} dx}{\int_{\mathbb{R}^n} e^{-f(x)} dx}$. Let x^* be the mode of ν , i.e., $x^* := \operatorname{argmax}_{x \in \mathbb{R}^n} \nu(x)$. Then one can show that \bar{x} and x^* are not too far from each other. Further, on average, a sample from ν is close to both \bar{x} and x^* . Formally, we have the following result.

Lemma 1. *Suppose ν is α -strongly log-concave. Then it holds that*

1. $\operatorname{Var}_\nu(x) := \mathbb{E}_\nu[\|x - \bar{x}\|_2^2] \leq \frac{n}{\alpha},$
2. $\mathbb{E}_\nu[\|x - x^*\|_2^2] \leq \frac{n}{\alpha},$
3. $\|\bar{x} - x^*\|_2^2 \leq \frac{n}{\alpha}.$

Example 2. *Consider the Gaussian distribution $\rho = \mathcal{N}(\nu, \Sigma)$ with mean $\nu \in \mathbb{R}^n$ covariance $\Sigma \in \mathbb{R}^{n \times n}$. From Example 1, we know that ρ is $\frac{1}{\lambda_{\max}(\Sigma)}$ -strongly log-concave. In this case, $\operatorname{Var}_\nu(x) = \operatorname{Tr}(\Sigma) = \sum_{i=1}^n \lambda_i(\Sigma) \leq n \cdot \lambda_{\max}(\Sigma) = \frac{n}{\alpha}.$*

See [DCWY19, Section 5.3] for an outline of the proof of Lemma 1. Lemma 1 suggests that “effectively” samples from ν lie at a distance of at most $\approx \sqrt{n/\alpha}$ from each other. $\sqrt{n/\alpha}$ is popularly known as the effective diameter of the distribution. Lemma 2 formalizes this idea.

For all $s > 0$, define $r(s) := 2 + 2 \cdot \max\{(n^{-1} \log s^{-1})^{\frac{1}{4}}, (n^{-1} \log s^{-1})^{\frac{1}{2}}\}$. Define $\mathcal{R}_s \subseteq \mathbb{R}^n$ as

$$\mathcal{R}_s := \mathbb{B}(x^*, r(s) \cdot \sqrt{n/\alpha}) \quad (2)$$

Lemma 2 ([DCWY19, Lemma 1]). *If ν is α -strongly log-concave on \mathbb{R}^n , then $\nu(\mathcal{R}_s) \geq 1 - s$.*

Note that, if $s = e^{-cn}$ then $r(s) \approx \sqrt{c}$. Therefore, by Lemma 2, at most $e^{-\Omega(n)}$ fraction of the samples from ν lie farther than $O(\sqrt{n/\alpha})$ distance from its x^* . Using this, one can show that the restriction of ν to \mathcal{R}_s is $(\frac{s}{1-s})$ -close in total variation distance to ν .² Thus, for a sufficiently small s , it suffices to restrict ν to the ball \mathcal{R}_s .

Interlude to optimization. Equivalently, one can define x^* , the mode of ν , as the minimizer of f (where $\nu \propto e^{-f}$), i.e.,

$$x^* = \operatorname{argmin}_{x \in \mathbb{R}^n} f(x).$$

Since ν is α -strongly log-concave, f is α -strongly convex. Assume that f is L -smooth. Therefore, we have $\alpha I \preceq \nabla^2 f(x) \preceq LI$. Define the *condition number* of f as $\kappa := L/\alpha$. As we will see in later

²To see this, note that $\nu_s(A) - \nu(A) = \frac{\nu(A \cap \mathcal{R}_s)}{\nu(\mathcal{R}_s)} - \nu(A) \geq \frac{\nu(A) - \nu(\mathcal{R}_s^c)}{\nu(\mathcal{R}_s)} - \nu(A) \geq \nu(A) - s - \nu(A) = -s$ and $\nu_s(A) - \nu(A) = \frac{\nu(A \cap \mathcal{R}_s)}{\nu(\mathcal{R}_s)} - \nu(A) \leq \frac{\nu(A)}{\nu(\mathcal{R}_s)} - \nu(A) \leq \frac{\nu(A)}{1-s} - \nu(A) \leq \frac{s}{1-s}.$

lectures, it is a theorem that gradient descent finds x^* using $\tilde{O}(\kappa)$ calls to first-order oracle for f [Vis21, Chapter 6]. Further, it is also a theorem that, accelerated gradient descent finds x^* using $\tilde{O}(\sqrt{\kappa})$ calls to first-order oracle for f [Vis21, Chapter 8] and this is optimal: any gradient-based algorithm requires at least $\Omega(\sqrt{\kappa})$ oracle-calls [Vis21, Exercise 8.4].

4 Isoperimetry with one-step overlap implies s -conductance

In this section, we show that if a Markov chain satisfies a “one-step overlap” property and its stationary distribution’s isoperimetry constant is not too large, then its s -conductance is not too small. Formally, we prove Theorem 2.

Theorem 2 (Isoperimetry with one-step overlap implies s -conductance [DCWY19, Lemma 2]). *Suppose ν is a α -strongly log-concave distribution on \mathbb{R}^n . Let P be a Markov chain with stationary distribution ν . For all $s \in (0, 1/2]$ and $\Delta_s \geq 0$, if P satisfies the following one-step overlap property (in the euclidean norm):*

$$\text{For all } x, y \in \mathcal{R}_s, \quad \text{if } \|x - y\|_2 \leq \Delta_s, \quad \text{then } \text{TV}(P_x, P_y) \leq \frac{3}{4}, \quad (3)$$

then P has s -conductance at least $\min\{1/16, \sqrt{\alpha} \cdot \Delta_s / 128\}$, i.e.,

$$\phi_s(P) \geq \min \left\{ \frac{1}{16}, \frac{\sqrt{\alpha} \cdot \Delta_s}{128} \right\}. \quad (4)$$

Recall that, from Theorem 1 we know that any strongly log-concave distribution has a small isoperimetry constant. Thus, any Markov chain P with the one-step overlap property and a α -strongly log-concave distribution, has s -conductance at least $\sqrt{\alpha} \cdot \Delta_s / 128$. Intuitively, the property requires that if two points are close together, then the next-step distributions on starting P from either point are also close to each other. It turns out several fundamental Markov chains have this property, including the Metropolis random walk (e.g., [LS93]) and MALA (e.g., [MV19]). We will prove this property for the Metropolis random walk in the next lecture. (Somewhat interestingly, it also holds for Markov chains like hit-and-run [Lov99] and Gibbs-sampler [LV20, NS20] which take large steps in the domain.)

Remark 1 (The constant in the one-step property). *The constant $3/4$ in Equation (3) is not crucial. In fact, for any constant $c > 0$, one can replace $3/4$ with $1 - c$ in Equation (3). This would only change other constants in Theorem 2.*

Combining Theorem 2 with a bound on the mixing-time in terms of s -conductance (from Lecture 6), we get the following corollary.

Corollary 1 (Bound on mixing time). *Let P be a Markov chain with stationary distribution ν . For all $s \in (0, 1/2]$ and $\Delta_s \geq 0$, if P satisfies the following one-step overlap property (in the euclidean norm):*

$$\text{For all } x, y \in \mathcal{R}_s, \quad \text{if } \|x - y\|_2 \leq \Delta_s, \quad \text{then } \text{TV}(P_x, P_y) \leq \frac{3}{4}, \quad (5)$$

then starting from an M -warm start P converges to a distribution ε -close in total variation distance to ν in $O(\phi_t^{-2} \cdot \log(M/\varepsilon)) = O(\alpha^{-1} \cdot \Delta_t^{-2} \cdot \log(M/\varepsilon))$ steps, where $t = \varepsilon/2M$.

4.1 Proof of Theorem 2

Fix any set $A \subseteq \mathbb{R}^n$ with $s < \nu(A) < 1 - s$. We will show that $C(A) \geq \min\{1/16, \sqrt{\alpha} \cdot \Delta_s/128\}$, where

$$C(A) := \frac{\Phi(A)}{\min\{\nu(A) - s, \nu(A^c) - s\}}.$$

Then, Theorem 2 follows by using Definition 1.

Given $A \subseteq \mathbb{R}^n$, define sets $A_1 \subseteq A$ and $A_2 \subseteq A^c$, containing points “deep inside” A and A^c respectively, as

$$A_1 := \{x \in A \cap \mathcal{R}_s \mid P_x(A^c) < 1/8\}, \quad (6)$$

$$A_2 := \{y \in A^c \cap \mathcal{R}_s \mid P_y(A) < 1/8\}. \quad (7)$$

Further, define $A_3 := \mathcal{R}_s \setminus A_1 \setminus A_2$ as the remaining portion of \mathcal{R}_s . Thus, $\{A_1, A_2, A_3\}$ partitions \mathcal{R}_s .

Intuitively, points in A_1 (resp. A_2) are unlikely to cross over to A^c (resp. A) in the next step. Therefore, if the measure of A_1 or A_2 is small (relative to $\nu(A)$ and $\nu(A^c)$), then the ergodic flow $\Phi(A)$ should be large. Lemma 3 formalizes this idea.

Lemma 3. *If $\nu(A_1) \leq \frac{1}{2} \cdot \nu(A \cap \mathcal{R}_s)$ or $\nu(A_2) \leq \frac{1}{2} \cdot \nu(A^c \cap \mathcal{R}_s)$, then $C(A) \geq \frac{1}{16}$.*

Proof. First, suppose that $\nu(A_1) \leq \frac{1}{2} \cdot \nu(A \cap \mathcal{R}_s)$. In this case we have

$$\begin{aligned} \Phi(A) &= \int_{x \in A} P_x(A^c) d\nu(x) \\ &\geq \int_{x \in (A \cap \mathcal{R}_s) \setminus A_1} P_x(A^c) d\nu(x) \\ &\stackrel{(6)}{\geq} \frac{1}{8} \cdot \nu((A \cap \mathcal{R}_s) \setminus A_1) \\ &\geq \frac{1}{16} \cdot \nu(A \cap \mathcal{R}_s) \quad \text{(Using that } \nu(A_1) \leq \frac{1}{2} \cdot \nu(A \cap \mathcal{R}_s)\text{)} \\ &= \frac{1}{16} \cdot (\nu(A) - \nu(A \cap \mathcal{R}_s^c)) \\ &\geq \frac{1}{16} \cdot (\nu(A) - \nu(\mathcal{R}_s^c)) \\ &\stackrel{\text{Lemma 2}}{\geq} \frac{1}{16} \cdot (\nu(A) - s) \end{aligned}$$

$$\geq \frac{1}{16} \cdot \min \{ \nu(A) - s, \nu(A^c) - s \}.$$

Thus, it follows that

$$C(A) = \frac{\Phi(A)}{\min \{ \nu(A) - s, \nu(A^c) - s \}} \geq \frac{1}{16}.$$

By symmetry, if $\nu(A_2) \leq \frac{1}{2} \cdot \nu(A^c \cap \mathcal{R}_s)$, then we have $C(A^c) \geq \frac{1}{16}$. Since ν is the stationary distribution of P , we have $C(A) = C(A^c)$. Thus, $C(A) \geq \frac{1}{16}$. \square

It remains to prove Equation (4), when $\nu(A_1) \geq \frac{1}{2} \cdot \nu(A \cap \mathcal{R}_s)$ and $\nu(A_2) \geq \frac{1}{2} \cdot \nu(A^c \cap \mathcal{R}_s)$.

Lemma 4. *If $\nu(A_1) \geq \frac{1}{2} \cdot \nu(A \cap \mathcal{R}_s)$ and $\nu(A_2) \geq \frac{1}{2} \cdot \nu(A^c \cap \mathcal{R}_s)$, then $\Phi(A) \geq \frac{1}{16} \cdot \nu(A_3)$.*

Proof.

$$\begin{aligned} \Phi(A) &= \frac{1}{2} \cdot (\Phi(A) + \Phi(A^c)) && \text{(Using that } \nu \text{ is the stationary distribution of } P\text{)} \\ &= \frac{1}{2} \cdot \left(\int_{x \in A} P_x(A^c) d\nu(v) + \int_{x \in A^c} P_x(A) d\nu(v) \right) \\ &\geq \frac{1}{2} \cdot \left(\int_{x \in (A \cap \mathcal{R}_s) \setminus A_1} P_x(A^c) d\nu(v) + \int_{x \in (A^c \cap \mathcal{R}_s) \setminus A_2} P_x(A) d\nu(v) \right) \\ &\stackrel{(6),(7)}{\geq} \frac{1}{16} \cdot \nu(A_3). \end{aligned} \tag{8}$$

\square

Towards proving Equation (4), we would like to lower bound $\nu(A_3)$ when $\nu(A_1)$ and $\nu(A_2)$ are not too small—they satisfy

$$\nu(A_1) \geq \frac{1}{2} \cdot \nu(A \cap \mathcal{R}_s) \quad \text{and} \quad \nu(A_2) \geq \frac{1}{2} \cdot \nu(A^c \cap \mathcal{R}_s). \tag{9}$$

This is where isoperimetry comes in. However, we cannot use Equation (1) with the measure ν and sets $A_1, A_2, A_3 \subseteq \mathbb{R}^n$, because A_1, A_2 , and A_3 do not partition \mathbb{R}^n . We bypass this by considering the restriction, ν_s , of ν to \mathcal{R}_s . Formally, let ν_s be the distribution satisfying that for all sets $B \subseteq \mathbb{R}^n$ that

$$\nu_s(B) := \frac{\nu(B \cap \mathcal{R}_s)}{\nu(\mathcal{R}_s)}. \tag{10}$$

Fact 1. *If ν is α -strongly log-concave and \mathcal{R}_s is a convex set, then ν_s is also α -strongly log-concave.*

Fact 1 allows us to use Equation (1) for ν_s with the partition $\mathcal{R}_s = A_1 \cup A_2 \cup A_3$.

$$\nu_s(A_3) \geq \frac{\sqrt{\alpha}}{4} \cdot d(A_1, A_2) \cdot \min \{\nu_s(A_1), \nu_s(A_2)\}$$

We can transfer this lower bound on $\nu_s(A_3)$ to a lower bound on $\nu(A_3)$ by scaling the above inequality by $\nu(\mathcal{R}_s) > 0$ and using Equation (10).

$$\begin{aligned} \nu(A_3) &\geq \frac{\sqrt{\alpha}}{4} \cdot d(A_1, A_2) \cdot \min \{\nu(A_1), \nu(A_2)\} \\ &\stackrel{(9)}{\geq} \frac{\sqrt{\alpha}}{8} \cdot d(A_1, A_2) \cdot \min \{\nu(A) - s, \nu(A^c) - s\}. \end{aligned} \quad (11)$$

Finally, we require a lower bound $d(A_1, A_2)$. This follows from the one-step overlap property.

Lemma 5. *It holds that $d(A_1, A_2) \geq \Delta_s$.*

Proof. For any $x \in A_1$ and $y \in A_2$, we have that

$$\begin{aligned} \text{TV}(P_x, P_y) &= \sup_{B \subseteq \mathbb{R}^n} |P_x(B) - P_y(B)| \\ &\geq P_x(A) - P_y(A) \\ &= 1 - P_x(A^c) - P_y(A) \\ &\geq 1 - \frac{1}{8} - \frac{1}{8} \\ &= \frac{3}{4}. \end{aligned}$$

Therefore, by the (contrapositive of) one-step overlap property (Equation (5)), it must hold that

$$\|x - y\|_2 > \Delta_s. \quad (12)$$

Since this holds for all $x \in A_1$ and $y \in A_2$, we have that

$$d(A_1, A_2) = \inf_{x \in A_1, y \in A_2} \|x - y\|_2 \stackrel{(12)}{\geq} \Delta_s.$$

(Notice that we loose the strict inequality, why?) □

Combining Lemma 5 with Equation (11), we get that, if Equation (9) holds, then

$$\nu(A_3) \geq \frac{\sqrt{\alpha} \cdot \Delta_s}{8} \cdot \min \{\nu(A) - s, \nu(A^c) - s\}.$$

Substituting this into Lemma 4, we get

$$\Phi(A) \geq \frac{\sqrt{\alpha} \cdot \Delta_s}{128} \cdot \min \{\nu(A) - s, \nu(A^c) - s\}.$$

Thus, $C(A) \geq \frac{1}{128} \cdot \sqrt{\alpha} \cdot \Delta_s$. Combining this with Lemma 3 completes the proof.

5 An aside on the Kullback-Leibler divergence

Let ρ and ν be probability distributions on some domain \mathcal{X} . The Kullback-Leibler (KL) divergence of ρ with respect to ν is defined as

$$H_\nu(\rho) := \int_{x \in \mathcal{X}} \rho(x) \cdot \log \frac{\rho(x)}{\nu(x)} dx.$$

KL divergence is also known as relative-entropy, due to its close connection to the Shanon entropy [Sha48]. KL divergence has several nice properties—Lemma 6 summarizes some of them.

Lemma 6. *The Kullback-Leibler (KL) divergence satisfies the following properties:*

1. For all distributions ρ and ν , $H_\nu(\rho) \geq 0$.
2. For all distributions ρ and ν , $H_\nu(\rho) \geq 0$ if and only if $\rho = \nu$.
3. For all distributions ρ and ν , if ρ has density $h = \frac{d\rho}{d\nu} : \mathcal{X} \rightarrow \mathbb{R}$ w.r.t. ν , then

$$H_\nu(\rho) = \mathbb{E}_\nu[h(x) \log h(x)].$$

Note that KL divergence is *not* a metric, in particular, it is not symmetric and does not satisfy the triangle inequality.

We are interested in KL divergence because of its relation to the total variation distance: For any distributions ρ and ν , it holds that

$$\text{TV}(\rho, \nu) \leq \sqrt{\frac{1}{2} \cdot H_\nu(\rho)}. \quad (\text{Pinsker's inequality})$$

In the next lecture, we will use KL divergence and Pinsker's inequality to show that Metropolis Random Walk (MRW) satisfies the one-step overlap property (in total variation distance).

References

- [Bal06] K. Ball. *Lecture*. Institut Henri Poincare, Paris, 2006.
- [DCWY19] Raaz Dwivedi, Yuansi Chen, Martin J. Wainwright, and Bin Yu. Log-concave sampling: Metropolis-hastings algorithms are fast. *J. Mach. Learn. Res.*, 20:183:1–183:42, 2019.
- [EL14] Ronen Eldan and Joseph Lehec. Bounding the norm of a log-concave vector via thin-shell estimates. In *Geometric Aspects of Functional Analysis*, pages 107–122. Springer, 2014.
- [Eld13] Ronen Eldan. Thin shell implies spectral gap up to polylog via a stochastic localization scheme. *Geometric and Functional Analysis*, 23(2):532–569, 2013.

- [FGP07] Bruno Fleury, Olivier Guédon, and Grigoris Paouris. A stability result for mean width of ℓ_p -centroid bodies. *Advances in Mathematics*, 214(2):865–877, 2007.
- [Fle10] Bruno Fleury. Concentration in a thin euclidean shell for log-concave measures. *Journal of Functional Analysis*, 259(4):832–841, 2010.
- [GM11] Olivier Guédon and Emanuel Milman. Interpolating thin-shell and sharp large-deviation estimates for isotropic log-concave measures. *Geometric and Functional Analysis*, 21(5):1043, 2011.
- [Kla07a] B Klartag. Power-law estimates for the central limit theorem for convex sets. *Journal of Functional Analysis*, 245(1):284–310, 2007.
- [Kla07b] Bo’az Klartag. A central limit theorem for convex sets. *Inventiones mathematicae*, 168(1):91–131, 2007.
- [KLS95] Ravi Kannan, László Lovász, and Miklós Simonovits. Isoperimetric problems for convex bodies and a localization lemma. *Discret. Comput. Geom.*, 13:541–559, 1995.
- [Lov99] László Lovász. Hit-and-run mixes fast. *Mathematical Programming*, 86(3):443–461, 1999.
- [LS93] László Lovász and Miklós Simonovits. Random walks in a convex body and an improved volume algorithm. *Random structures & algorithms*, 4(4):359–412, 1993.
- [LV07] László Lovász and Santosh Vempala. The geometry of logconcave functions and sampling algorithms. *Random Structures & Algorithms*, 30(3):307–358, 2007.
- [LV17a] Yin Tat Lee and Santosh Srinivas Vempala. Eldan’s stochastic localization and the kls hyperplane conjecture: An improved lower bound for expansion. In *2017 IEEE 58th Annual Symposium on Foundations of Computer Science (FOCS)*, pages 998–1007. IEEE, 2017.
- [LV17b] Yin Tat Lee and Santosh Srinivas Vempala. Eldan’s stochastic localization and the KLS hyperplane conjecture: An improved lower bound for expansion. In *FOCS*, pages 998–1007. IEEE Computer Society, 2017.
- [LV18] Yin Tat Lee and Santosh S Vempala. The kannan-lovász-simonovits conjecture. *arXiv preprint arXiv:1807.03465*, 2018.
- [LV20] Aditi Laddha and Santosh Vempala. Convergence of gibbs sampling: Coordinate hit-and-run mixes fast. *arXiv preprint arXiv:2009.11338*, 2020.

- [MV19] Oren Mangoubi and Nisheeth K. Vishnoi. Nonconvex sampling with the metropolis-adjusted langevin algorithm. In *COLT*, volume 99 of *Proceedings of Machine Learning Research*, pages 2259–2293. PMLR, 2019.
- [NS20] Hariharan Narayanan and Piyush Srivastava. On the mixing time of coordinate hit-and-run. *CoRR*, abs/2009.14004, 2020.
- [Sha48] Claude E Shannon. A mathematical theory of communication. *The Bell system technical journal*, 27(3):379–423, 1948.
- [Vem05] Santosh Vempala. Geometric random walks: a survey. *Combinatorial and computational geometry*, 52(573-612):2, 2005.
- [Vis21] Nisheeth K. Vishnoi. *Algorithms for Convex Optimization*. Cambridge University Pres, 2021.
- [YTL19] Santosh Vempala Yin Tat Lee. *Techniques in Optimization and Sampling*. 2019.